A Quick Guide to Networking Software
by J. Sanguino
5th Edition: February 2020

# Mission Briefing

Welcome to this quick guide in networking programming.
You will be given a username and password to access any of the RC
lab computers. They are connected to Internet and running Linux.

Your mission, should you decide to accept it, is to complete the
tasks that will be presented as you move along the guide. They
involve the development of programs that communicate through the
Internet.
The tools that you will be using are the basis for the development
of network applications over the Internet (web browsers and
servers, email, peer-to-peer, remote logins, file transfers …).

The kind of network applications you will be able to develop, on
your own, at the end of this guide, will only be bounded by your
imagination.

As always, should you or any team member be caught in thrall of
network programming, the author would disavow any knowledge of your
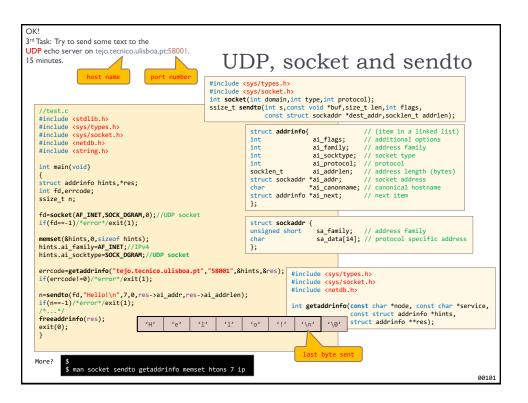actions.

00001

1

## Slide 1

Login:      `alunos`

Password:   `alunos`

## Slide 2

Welcome, you are inside now.
1st Task:  Get the host name!
You have 10 minutes.

# gethostname

```
#include <unistd.h>
int gethostname(char *name,size_t len);
```

```c
//test.c
#include <stdio.h>
#include <stdlib.h>
#include <unistd.h>
#include <string.h>
#include <errno.h>
extern int errno;

int main(void)
{
char buffer[128];

if(gethostname(buffer,128)==-1)
    fprintf(stderr,"error: %s\n",strerror(errno));

else printf("host name: %s\n",buffer);
exit(0);
}
```

```makefile
# makefile

test: test.c
⟵⟶ gcc test.c -o test
 tab
```

running on tejo.tecnico.ulisboa.pt

```
$ make
gcc test.c -o test
$ ./test
host name: tejo.tecnico.ulisboa.pt
$
```

```
#include <string.h>
char *strerror(int errnum);
```

More?
```
$
$ man gethostname strerror
```

## Slide 1 (00100)

Good! Move on!
2nd Task: Now that you have a name, get the IP address.
15 minutes.

# getaddrinfo

```c
#include <sys/types.h>
#include <sys/socket.h>
#include <netdb.h>

int getaddrinfo(const char *node, const char *service,
                const struct addrinfo *hints,
                struct addrinfo **res);
void freeaddrinfo(struct addrinfo *res);
const char *gai_strerror(int errcode);
```

```c
#include <arpa/inet.h>

const char *inet_ntop(int af,
    const void *src, char *dst,
    socklen_t size);
```

```c
// test.c
#include <stdio.h>
#include <stdlib.h>
#include <sys/types.h>
#include <sys/socket.h>
#include <netdb.h>
#include <arpa/inet.h>
#include <string.h>

int main(void)
{
struct addrinfo hints,*res,*p;
int errcode;
char buffer[INET_ADDRSTRLEN];
struct in_addr *addr;

memset(&hints,0,sizeof hints);
hints.ai_family=AF_INET;//IPv4
hints.ai_socktype=SOCK_DGRAM;
hints.ai_flags=AI_CANONNAME;

if((errcode=getaddrinfo("tejo.tecnico.ulisboa.pt",NULL,&hints,&res))!=0)
    fprintf(stderr,"error: getaddrinfo: %s\n",gai_strerror(errcode));
else{
    printf("canonical hostname: %s\n",res->ai_canonname);
    for(p=res;p!=NULL;p=p->ai_next){
        addr=&((struct sockaddr_in *)p->ai_addr)->sin_addr;
        printf("internet address: %s (%08lX)\n",
            inet_ntop(p->ai_family,addr,buffer,sizeof buffer),
            (long unsigned int)ntohl(addr->s_addr));
    }
    freeaddrinfo(res);
    }
exit(0);
}
```

```c
#include <string.h>
void *memset(void *s,int c,size_t n);
```

```c
struct addrinfo {            // (item in a linked list)
int            ai_flags;     // additional options
int            ai_family;    // address family
int            ai_socktype;  // socket type
int            ai_protocol;  // protocol
socklen_t      ai_addrlen;   // address length (bytes)
struct sockaddr *ai_addr;    // socket address
char           *ai_canonname; // canonical hostname
struct addrinfo *ai_next;    // next item
};
```

```c
struct sockaddr_in {
sa_family_t   sin_family;// address family: AF_INET
u_int16_t     sin_port;  // port in (16 bits)
struct in_addr sin_addr; // internet address
};
```
network byte order

```c
struct in_addr{
uint32_t   s_addr; // 32 bits
};
```

```c
#include <arpa/inet.h>
uint32_t ntohl(uint32_t netlong);
        (network to host long)
```

Long (32 bits) 0x76543210
Little endian system    Network byte order
  ADDR   0x10             ADDR   0x76
  ADDR+1 0x32             ADDR+1 0x54
  ADDR+2 0x54             ADDR+2 0x32
  ADDR+3 0x76             ADDR+3 0x10

Big Endian

0xC1==193
0x88==136
0x8A==138
0x8E==142

```
$ make
gcc test.c -o test
$ ./test
canonical hostname: tejo.tecnico.ulisboa.pt
internet address: 193.136.138.142 (C1888A8E)
$
```

More?

```
$
$ man getaddrinfo inet_ntop memset ntohl 7 ip
```

## Slide 2 (00101)

OK!
3rd Task: Try to send some text to the
UDP echo server on tejo.tecnico.ulisboa.pt:58001.
15 minutes.

host name    port number

# UDP, socket and sendto

```c
#include <sys/types.h>
#include <sys/socket.h>
int socket(int domain,int type,int protocol);
ssize_t sendto(int s,const void *buf,size_t len,int flags,
               const struct sockaddr *dest_addr,socklen_t addrlen);
```

```c
//test.c
#include <stdlib.h>
#include <sys/types.h>
#include <sys/socket.h>
#include <netdb.h>
#include <string.h>

int main(void)
{
struct addrinfo hints,*res;
int fd,errcode;
ssize_t n;

fd=socket(AF_INET,SOCK_DGRAM,0);//UDP socket
if(fd==-1)/*error*/exit(1);

memset(&hints,0,sizeof hints);
hints.ai_family=AF_INET;//IPv4
hints.ai_socktype=SOCK_DGRAM;//UDP socket

errcode=getaddrinfo("tejo.tecnico.ulisboa.pt","58001",&hints,&res);
if(errcode!=0)/*error*/exit(1);

n=sendto(fd,"Hello!\n",7,0,res->ai_addr,res->ai_addrlen);
if(n==-1)/*error*/exit(1);
/*...*/
freeaddrinfo(res);
exit(0);
}
```

```c
struct addrinfo{            // (item in a linked list)
int            ai_flags;    // additional options
int            ai_family;   // address family
int            ai_socktype; // socket type
int            ai_protocol; // protocol
socklen_t      ai_addrlen;  // address length (bytes)
struct sockaddr *ai_addr;   // socket address
char           *ai_canonname; // canonical hostname
struct addrinfo *ai_next;   // next item
};
```

```c
struct sockaddr {
unsigned short  sa_family;  // address family
char            sa_data[14]; // protocol specific address
};
```

```c
#include <sys/types.h>
#include <sys/socket.h>
#include <netdb.h>

int getaddrinfo(const char *node, const char *service,
                const struct addrinfo *hints,
                struct addrinfo **res);
```

| 'H' | 'e' | 'l' | 'l' | 'o' | '!' | '\n' | '\0' |

last byte sent

More?

```
$
$ man socket sendto getaddrinfo memset htons 7 ip
```

## Slide 1

4th Task: Now, receive the echo from the UDP echo server.
20 minutes.

# UDP and recvfrom

```c
//test.c
#include <stdio.h>
#include <stdlib.h>
#include <sys/types.h>
#include <sys/socket.h>
#include <netdb.h>
#include <string.h>
#include <unistd.h>

int main(void)
{
int fd;
struct sockaddr addr;
socklen_t addrlen;
ssize_t n;
char buffer[128+1];

/*...*///see previous task code

addrlen=sizeof(addr);
n=recvfrom(fd,buffer,128,0,&addr,&addrlen);
if(n==-1)/*error*/exit(1);
buffer[n] = '\0';
printf("echo: %s\n", buffer);

close(fd);
exit(0);
}
```

```c
#include <sys/types.h>
#include <sys/socket.h>
ssize_t  recvfrom(int s,void *buf,size_t len,int flags,
                  struct sockaddr *from,socklen_t *fromlen);
```

```
$ make
gcc test.c -o test
$ ./test
echo: Hello!
$
```

input/output argument

**Question 2**: How do you know the message you received came from the UDP echo server on tejo:58001.

**Question 3**: Which port number is your UDP client listening to when it is waiting for the echo reply?

**Question 4**: How many bytes do you expect to receive from recvfrom?

**Question 1**: What happens if the messages do not arrive at the destination? Try specifying a wrong port number for the destination echo server. Did you get an error message?

**Question 5**: Do you expect buffer content to be a NULL terminated string?

More?
```
$
$ man recvfrom
```

00110

## Slide 2

# Answers

**Answer to question 1:** No message will be received back at the client and it will block in recvfrom. No error will be detected unless timeouts are used.
You are using UDP. There are no guarantees that the messages will be delivered at the destination, and the order by which they are delivered may not be the same in which they were transmitted.

**Answer to question 2:** You have to check the recvfrom addr output argument. See, in the next slide, how to use getnameinfo for that purpose.
If you only want to receive messages from a specific address, then use send and recv. Find out more on manual page 2 (man 2 send recv).

**Answer to question 3:** The system assigned some unused port in the range 1024 through 5000 when you first called sendto and this is the port recvfrom is listening to.
If you want to use a specific port number you have to use bind. More on that later.

**Answer to question 4:** In this particular case, you should expect to receive 7 bytes (see sendto in previous slide).

**Answer to question 5:** In this particular case, you should not expect buffer to be NULL terminated. See sendto in previous slide and notice that the '\0' was not sent.

**Question 2**: How do you know the message you received came from the UDP echo server on tejo:58001.

**Question 3**: Which port number is your UDP client listening to when it is waiting for the echo reply?

**Question 4**: How many bytes do you expect to receive from recvfrom?

**Question 1**: What happens if the messages do not arrive at the destination? Try specifying a wrong port number for the destination echo server. Did you get an error message?

**Question 5**: Do you expect buffer content to be a NULL terminated string?

00111

4

# getnameinfo

```c
#include <sys/socket.h>
#include <netdb.h>
int getnameinfo(const struct sockaddr *addr, socklen_t addrlen,
                char *host, socklen_t hostlen,
                char *serv, socklen_t servlen, int flags);
```

```c
//test.c
#include <stdio.h>
#include <netdb.h>
#include <sys/socket.h>
/* ... */

int main(void)
{
int fd;
struct sockaddr addr;
socklen_t addrlen;
ssize_t n;
char buffer[128];
int errcode;
char host[NI_MAXHOST],service[NI_MAXSERV];//consts in <netdb.h>

/*...*/// see previous task code

addrlen=sizeof(addr);
n=recvfrom(fd,buffer,128,0,&addr,&addrlen);
if(n==-1)/*error*/exit(1);
/*...*/

if((errcode=getnameinfo(&addr,addrlen,host,sizeof(host),service,sizeof(service),0))!=0)
    fprintf(stderr,"error: getnameinfo: %s\n",gai_strerror(errcode));
else
    printf("sent by [%s:%s]\n",host,service);

exit(0);
}
```

```
$ make
gcc test.c -o test
$ ./test
echo: Hello!
sent by [tejo.tecnico.ulisboa.pt:58001]
$
```

More?

```
$
$ man getnameinfo
```

01000

---

OK. Now let's move from UDP to TCP.
TCP is connection-oriented.
6th Task: Connect to the TCP echo server on tejo.tecnico.ulisboa:58001.
10 minutes.

# TCP, socket and connect

```c
#include <sys/types.h>
#include <sys/socket.h>
int connect(int  sockfd,const struct sockaddr *serv_addr,
            socklen_t addrlen);
```

```c
//test.c
#include <stdlib.h>
#include <sys/types.h>
#include <sys/socket.h>
#include <netdb.h>
#include <string.h>

int main(void)
{
struct addrinfo hints,*res;
int fd,n;

fd=socket(AF_INET,SOCK_STREAM,0);//TCP socket
if(fd==-1)exit(1);//error

memset(&hints,0,sizeof hints);
hints.ai_family=AF_INET;//IPv4
hints.ai_socktype=SOCK_STREAM;//TCP socket

n=getaddrinfo("tejo.tecnico.ulisboa.pt","58001",&hints,&res);
if(n!=0)/*error*/exit(1);

n=connect(fd,res->ai_addr,res->ai_addrlen);
if(n==-1)/*error*/exit(1);

/*...*/
}
```

**Question 6**: Did you notice that the host name and port number are the same as before?

**Question 7**: What do you expect to happen if you type the wrong host name or port number?

Alternative:
errcode=getaddrinfo("193.136.138.142","58001",&hints,&res);

More?

```
$
$ man connect
```

01001

5

# Answers

**Answer to question 6:** There is no problem in having two servers on the same port number as long as they are using different protocols. In this case, one is using UDP and the other TCP.

**Question 7**: What do you expect to happen if you type the wrong host name or port number?

**Answer to question 7:** If you type the wrong host name, getaddrinfo would give you an error, unless you type a name that also exists. If you type the wrong port number, connect would give you an error, unless there is a TCP server listening on that port.

01010

---

7th Task: Send some text over the connection you have just established and read the response.
10 minutes.

# TCP, write and read

```c
//test.c
#include <stdio.h>
#include <unistd.h>
#include <string.h>
/* ... */
int main(void)
{
int fd;
ssize_t nbytes,nleft,nwritten,nread;
char *ptr,buffer[128+1];
/*...*////see previous task code
ptr=strcpy(buffer,"Hello!\n");
nbytes=7;
nleft=nbytes;
while(nleft>0){nwritten=write(fd,ptr,nleft);
            if(nwritten<=0)/*error*/exit(1);
            nleft-=nwritten;
            ptr+=nwritten;}
nleft=nbytes; ptr=buffer;
while(nleft>0){nread=read(fd,ptr,nleft);
            if(nread==-1)/*error*/exit(1);
            else if(nread==0)break;//closed by peer
            nleft-=nread;
            ptr+=nread;}
nread=nbytes-nleft;

buffer[nread] = '\0';
printf("echo: %s\n", buffer);
close(fd);
exit(0);
}
```

```c
#include <unistd.h>
ssize_t write(int fd,const void *buf,size_t count);
ssize_t read(int fd,void *buf,size_t count);
```

*also used to write and read to/from files*

```
$ make
gcc test.c -o test
$ ./test
echo: Hello!
$
```

**Question 8**: Did you notice that you may have to call write and read more than once?

**Question 9**: What do you expect to happen if your messages do not arrive at the destination?

More?

```
$
$ man 2 write read
```

01011

6

# Answers

**Answer to question 8:** There is no guarantee that `write` would send all the bytes you requested when you called it. Transport layer buffers may be full. However, `write` returns the number of bytes that were sent (accepted by the transport layer). So, you just have to use this information to make sure everything is sent.

You may also have to call `read` more that once, since `read` would return as soon as data is available at the socket. It may happen that, when `read` returns, there was still data to arrive. Since `read` returns the number of bytes read from the socket, you just have to use this information to make sure nothing is missing.

**Answer to question 9:** If the transport layer can not deliver your messages to the destination, the connection will be lost. In some circumstances, this may take a few minutes due to timeouts. If your process is blocked in a `read` when the connection is lost, then `read` would return -1 and `errno` would be set to the appropriate error.
If you call `write` on a lost connection, `write` would return -1, `errno` will be set to `EPIPE`, but the system would raise a `SIGPIPE` signal and, by default, that would kill your process. See the next slide for a way to deal with the `SIGPIPE` signal.
Note however that, if the connection is closed, by the peer process, in an orderly fashion, while `read` is blocking your process, then `read` would return 0, as a sign of `EOF` (end-of-file).
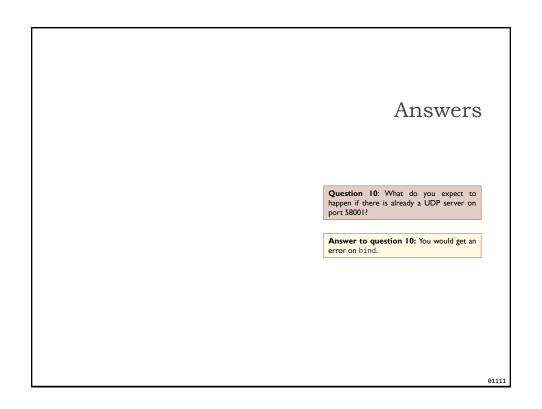
**Question 8**: Did you notice that you may have to call `write` and `read` more than once?

**Question 9**: What do you expect to happen if your messages do not arrive at the destination?
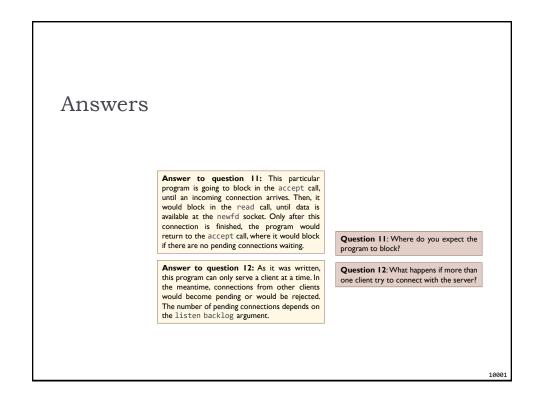
01100

---

Be careful. If the connection is lost and you write to the socket, the system will raise a SIGPIPE signal and, by default, this will kill your process.
8th Task: Protect the application against SIGPIPE signals.
5 minutes.

# TCP and the SIGPIPE signal

```c
#include <signal.h>
int sigaction(int signum, const struct sigaction *act,
              struct sigaction *oldact);
```

```c
//test.c
#include <signal.h>

/*...*/

int main(void)
{
/*...*/
struct sigaction act;

memset(&act,0,sizeof act);
act.sa_handler=SIG_IGN;

if(sigaction(SIGPIPE,&act,NULL)==-1)/*error*/exit(1);

/*...*/

}
```

```
                                                More?
$
$ man sigaction 7 signal
```

From now on, the `SIGPIPE` signal will be ignored.

Now, if the connection is lost and you write to the socket, the `write` will return -1 and `errno` will be set to `EPIPE`.

01101

7

Let's move from clients to servers.
Servers have well-known ports.
9th Task: Write a UDP echo server and run it on port 58001.
15 minutes.

# UDP server and bind

```
#include <stdlib.h>
#include <sys/types.h>
#include <sys/socket.h>
#include <netdb.h>
#include <string.h>

int main(void)
{
struct addrinfo hints,*res;
int fd,errcode;
struct sockaddr addr;
socklen_t addrlen;
ssize_t n,nread;
char buffer[128];

if((fd=socket(AF_INET,SOCK_DGRAM,0))==-1)exit(1);//error

memset(&hints,0,sizeof hints);
hints.ai_family=AF_INET;//IPv4
hints.ai_socktype=SOCK_DGRAM;//UDP socket
hints.ai_flags=AI_PASSIVE;
if((errcode=getaddrinfo(NULL,"58001",&hints,&res))!=0)/*error*/exit(1);

if((bind(fd,res->ai_addr,res->ai_addrlen)==-1)/*error*/exit(1);

while(1){addrlen=sizeof(addr);
        nread=recvfrom(fd,buffer,128,0, &addr,&addrlen);
        if(nread==-1)/*error*/exit(1);
        n=sendto(fd,buffer,nread,0,&addr,addrlen);
        if(n==-1)/*error*/exit(1);
        }
//freeaddrinfo(res);
//close(fd);
//exit(0);
}
```

well-known port number

```
#include <sys/types.h>
#include <sys/socket.h>
int bind(int sockfd,const struct sockaddr *my_addr,
         socklen_t addrlen);
```

Use bind to register the server well known address (and port) with the system.

More?

```
$
$ man 2 bind
```

**Question 10**: What do you expect to happen if there is already a UDP server on port 58001?

**Note**: You can also use bind to register the address (and port) in clients. In that case, if you set the port number to 0, the system assigns some unused port in the range 1024 through 5000.

Send only the bytes you read.

01110

# Answers

**Question 10**: What do you expect to happen if there is already a UDP server on port 58001?

**Answer to question 10:** You would get an error on bind.

01111

8

Now, do the same, but with TCP.
10th Task: Write a TCP echo server and run it also on port 58001.
20 minutes.

# TCP server, bind, listen and accept

```c
#include <stdlib.h>
#include <sys/types.h>
#include <sys/socket.h>
#include <netdb.h>
#include <string.h>
#include <unistd.h>

int main(void)
{
struct addrinfo hints,*res;
int fd,newfd,errcode;        ssize_t n,nw;
struct sockaddr addr;     socklen_t addrlen;
char *ptr,buffer[128];

if((fd=socket(AF_INET,SOCK_STREAM,0))==-1)exit(1);//error

memset(&hints,0,sizeof hints);
hints.ai_family=AF_INET;//IPv4
hints.ai_socktype=SOCK_STREAM;//TCP socket
hints.ai_flags=AI_PASSIVE;
if((errcode=getaddrinfo(NULL,"58001",&hints,&res))!=0)/*error*/exit(1);

if(bind(fd,res->ai_addr,res->ai_addrlen)==-1)/*error*/exit(1);
if(listen(fd,5)==-1)/*error*/exit(1);

while(1){addrlen=sizeof(addr);
        if((newfd=accept(fd,&addr,&addrlen))==-1)
            /*error*/exit(1);
        while((n=read(newfd,buffer,128))!=0){if(n==-1)/*error*/exit(1);
            ptr=&buffer[0];
            while(n>0){if((nw=write(newfd,ptr,n))<=0)/*error*/exit(1);
                    n-=nw; ptr+=nw;}
            }
        close(newfd);
        }
//freeaddrinfo(res);close(fd);exit(0);
}
```

```c
#include <sys/types.h>
#include <sys/socket.h>
int bind(int sockfd,const struct sockaddr *my_addr,
            socklen_t addrlen);
int listen(int sockfd,int backlog);
int accept(int sockfd,struct sockaddr *addr,
            socklen_t *addrlen);
```

Use bind to register the server well known address (and port) with the system.

Use listen to instruct the kernel to accept incoming connection requests for this socket. The backlog argument defines the maximum length the queue of pending connections may grow to.

Use accept to extract the first connection request on the queue of pending connections. Returns a socket associated with the new connection.

address of the connected peer process

**Question 11**: Where do you expect the program to block?

**Question 12**: What happens if more than one client try to connect with the server?

**Note**: Do not forget to protect your application against the SIGPIPE signal.

More?

```
$
$ man 2 bind listen accept 7 tcp
```

10000

---

# Answers

**Answer to question 11:** This particular program is going to block in the accept call, until an incoming connection arrives. Then, it would block in the read call, until data is available at the newfd socket. Only after this connection is finished, the program would return to the accept call, where it would block if there are no pending connections waiting.

**Answer to question 12:** As it was written, this program can only serve a client at a time. In the meantime, connections from other clients would become pending or would be rejected. The number of pending connections depends on the listen backlog argument.

**Question 11**: Where do you expect the program to block?

**Question 12**: What happens if more than one client try to connect with the server?

10001

# Slide 1

If you are already serving a client, send "busy\n" to new incoming clients.
11th Task: Change the previous code to do that.
15 minutes.

## select

```c
#include <sys/time.h>
#include <sys/types.h>
#include <unistd.h>
/* ... */
#define max(A,B) ((A)>=(B)?(A):(B))

int main(void)
{
int fd,newfd,afd=0;
fd_set rfds;
enum {idle,busy} state;
int maxfd,counter;
/*...*/
/*fd=socket(…);bind(fd,…);listen(fd,…);*/
state=idle;
while(1){FD_ZERO(&rfds);
        switch(state){
            case idle: FD_SET(fd,&rfds);maxfd=fd; break;
            case busy: FD_SET(fd,&rfds);FD_SET(afd,&rfds);maxfd=max(fd,afd); break;
            }//switch(state)

        counter=select(maxfd+1,&rfds,(fd_set*)NULL,(fd_set*)NULL,(struct timeval *)NULL);
        if(counter<=0)/*error*/exit(1);

        for(;counter;--counter)
            switch(state){
                case idle: if(FD_ISSET(fd,&rfds)){FD_CLR(fd,&rfds);
                            addrlen=sizeof(addr);
                            if((newfd=accept(fd,&addr,&addrlen))==-1)/*error*/exit(1);
                            afd=newfd;state=busy;}
                            break;
                case busy: if(FD_ISSET(fd,&rfds)){FD_CLR(fd,&rfds);
                            addrlen=sizeof(addr);
                            if((newfd=accept(fd,&addr,&addrlen))==-1)/*error*/exit(1);
                            /* ... write "busy\n" in newfd */
                            close(newfd);}
                           else if(FD_ISSET(afd,&rfds)){FD_CLR(afd,&rfds);
                            if((n=read(afd,buffer,128))!=0)
                                {if(n==-1)/*error*/exit(1);
                                 /* ... write buffer in afd */}
                            else{close(afd);state=idle;}//connection closed by peer
                            }
                           break;
                }//switch(state)
        }//while(1)
/*close(fd);exit(0);*/
}//main
```

```c
#include <sys/time.h>
#include <sys/types.h>
#include <unistd.h>
int select(int n,fd_set *readfds,fd_set *writefds,
           fd_set *exceptfds,struct timeval *timeout);
FD_CLR(int fd,fd_set *set);
FD_ISSET(int fd,fd_set *set);
FD_SET(int fd,fd_set *set);
FD_ZERO(fd_set *set);
```

Returns the number of file descriptors ready.

Blocks until one of the file descriptors, previously set in rfds, are ready to by read.

fd is ready

fd is ready

afd is ready

**Question 13**: And now, where do you expect the program to block?

More?

```
$
$ man 2 select
```

10010

# Slide 2

## Answers

**Answer to question 13:** This program is only going to block in the select call. It would not block neither in the accept call, neither in the read call, since those are only executed when their sockets are ready to be read (and so they have no reason to block).

**Question 13**: And now, where do you expect the program to block?

10011

## Slide 10100

12th Task: Make your server a concurrent server.
15 minutes.

# fork

```c
#include <stdlib.h>
#include <sys/types.h>
#include <sys/socket.h>
#include <netdb.h>
#include <unistd.h>
#include <string.h>
#include <signal.h>
#include <errno.h>
extern int errno;

int main(void)
{
struct sigaction act;
struct addrinfo hints,*res;
int fd, newfd,ret;          ssize_t n,nw;
struct sockaddr addr;    socklen_t addrlen;
char *ptr,buffer[128];
pid_t pid;
memset(&act,0,sizeof act);
act.sa_handler=SIG_IGN;
if(sigaction(SICHLD,&act,NULL)==-1)/*error*/exit(1);

if((fd=socket(AF_INET,SOCK_STREAM,0))==-1)exit(1);//error
memset(&hints,0,sizeof hints);
hints.ai_family=AF_INET;//IPv4
hints.ai_socktype=SOCK_STREAM;//TCP socket
hints.ai_flags=AI_PASSIVE;
if((ret=getaddrinfo(NULL,"58001",&hints,&res))!=0)/*error*/exit(1);
if(bind(fd,res->ai_addr,res->ai_addrlen)==-1)/*error*/exit(1);
if(listen(fd,5)==-1)/*error*/exit(1);
freeaddrinfo(res);//frees the memory allocated by getaddrinfo (no longer needed)

while(1){addrlen=sizeof(addr);
        do newfd=accept(fd,&addr,&addrlen);//wait for a connection
        while(newfd==-1&&errno==EINTR);
        if(newfd==-1)/*error*/exit(1);

        if((pid=fork())==-1)/*error*/exit(1);
        else if(pid==0)//child process
        {close(fd);
            while((n=read(newfd,buffer,128))!=0){if(n==-1)/*error*/exit(1);
                ptr=&buffer[0];
                while(n>0){if((nw=write(newfd,ptr,n))<=0)/*error*/exit(1);
                        n-=nw; ptr+=nw;}
            }
            close(newfd); exit(0);}
        //parent process
        do ret=close(newfd);while(ret==-1&&errno==EINTR);
        if(ret==-1)/*error*/exit(1);
        }
/*close(fd);exit(0);*/}
```

Use fork to create a new process for each new connection.

```c
#include <unistd.h>
pid_t fork(void);
```

Avoid zombies when child processes die.

**Note**: Do not forget to protect the child process against the SIGPIPE signal.

Create a child process for each new connection.

Parent process may be interrupted by SIG_CHLD signal (child process death).

child process

More?

```
$
$ man fork
```

## Slide 10101

# Further Reading

*Unix Network Programming: Networking APIs: Sockets and XTI (Volume 1), 2nd ed.*, W. Richard Stevens, 1998, Prentice-Hall PTR, ISBN 013490012X.

*Unix Network Programming: Networking APIs: The Sockets Networking API (Volume 1), 3rd ed.*, W. Richard Stevens, Bill Fenner, Andrew M. Rudoff, 2003, Addison-Wesley Professional, ISBN 0131411551.

sanguino@lx.it.pt